

ROGUE — Video Script

Your LLM is being jailbroken. ROGUE finds out before your users do. — *The red-team that never sleeps.*

358

ATTACKS TRACKED

8,321

TRIALS JUDGED

6

DEPLOYMENTS TESTED

19

OPEN-WEB SOURCES

5/5

BRIGHT DATA PRODUCTS

Target runtime 4:35 (under 5:00 cap) · Format: TIME – ON SCREEN – VOICEOVER · Read
~150 wpm · Live: rogue-eosin.vercel.app

1 Cold open — the hook

0:00–0:18

ON SCREEN: Black. The hero word-rotator alone, cycling fast: *jailbroken* → *prompt-injected* → *role-played* → *escalated*. Red words flicker against black.

VOICEOVER

"Right now, somebody on Reddit, on GitHub, on arXiv, just published a brand-new way to jailbreak a model exactly like the one you shipped to production. You won't read it. Your users will. **ROGUE reads it first.**"

(Beat. Cut to the live hero.)

2 What it is — the hero

0:18–0:48

ON SCREEN: Scroll-stop on the live hero. Both pills visible — *live* · *streaming the open web and powered by Bright Data* · *5/5 products* · *cost-optimized*. Slow-pan the stat trio counting up.

VOICEOVER

"This is ROGUE — *Real-time Open-web Generation of jailbreak Updates and Evaluation*. It's a continuous red-team for production LLMs that never sleeps. It harvests every new attack the moment it's posted, fires it at *your* exact deployment, and tells you which of your guardrails just broke."

VOICEOVER (AS NUMBERS LAND)

"Not a benchmark. Live data: **358 real attacks** tracked, **eight-thousand-three-hundred trials** judged, across **six** customer-style deployments. Every dot you'll see traces back to an attack a real person published — reproduced against a real config, graded by a real model."

3 The Bright Data engine

0:48–1:18

ON SCREEN: Scroll to the **Sources Marquee** — 19 source logos streaming past (Reddit, X, GitHub, Hugging Face, arXiv...), tagged by Bright Data product.

VOICEOVER

"The open web doesn't *want* to be read at scale — rate limits, bot detection, JavaScript walls, geo-blocks. That's the wall ROGUE had to get through, and that's where Bright Data comes in. ROGUE runs on **all five** Bright Data products at once — the MCP Server, Web Scraper API, SERP API, Web Unlocker, and Scraping Browser — fanning out across **nineteen sources**, from Reddit threads to arXiv preprints to raw GitHub jailbreak dumps. New attack on the web; new row in the matrix, within minutes."

4 How it thinks — the pipeline

1:18–1:52

ON SCREEN: The "How ROGUE thinks" 3-step section: *Harvest* → *Reproduce* → *Defend*. Let the animated arrows draw.

VOICEOVER

"Three loops, one outcome. **Harvest** — stream the latest jailbreaks off the open web. **Reproduce** — and this is the hard part — replay each one against a panel of deployment configs: your model, your system prompt, your tools. An attacker model *iterates*, the way a real adversary would. **Defend** — a separate judge model grades every single trial, attaches a ninety-five-percent confidence interval, and ships the result. No synthetic examples. No cherry-picking. Just: did it break, and how sure are we."

5 The matrix — the money shot

1:52–2:35

ON SCREEN: Click "See what's breaching → /matrix." Land on the breach matrix — families × deployments, green-to-red heatmap. Hover a glowing **red** cell, click it; the drawer slides out with the exact payload + breach trail + CIs. Hit ► **replay**.

VOICEOVER

"This is the matrix. Attack families down the side, your deployments across the top. Green is held. **Red is broken**. And you don't have to trust me — click any red cell..."

VOICEOVER (AFTER CLICK)

"...and ROGUE hands you the receipt: the *exact prompt* that cracked it, the model's own response, the judge's rationale, and the confidence interval. Press play, and you watch the breach happen turn by turn. This is the artifact a security team actually argues from — evidence, not vibes."

6 Originality — the bandit + stress tests

2:35–3:15

ON SCREEN: Scroll to the **Augmentation Showcase** cards, then the interactive **Augmentation Lab** — toggle persona / escalation / mutation and watch the estimated breach rate stack up.

 VOICEOVER

"Finding attacks is easy. Finding them *cheaply* is the trick. ROGUE runs a multi-armed bandit over **thirty-nine** harvest sources — it learns which corners of the web yield the most *novel* attacks per dollar and reallocates budget automatically. The top source is currently returning over **seven thousand new attacks for every dollar** spent."

 VOICEOVER

"Then it doesn't just replay an attack — it *stress-tests* your defense five ways: wraps it in a persuasion persona, escalates it across turns, mutates the wording to dodge pattern-matching, and measures how stubborn each config really is. Toggle them right here and watch your breach rate climb. That's the difference between *did this one string break you* and *is this class of attack going to break you*."

7 The differentiator — ROGUE is an MCP server

3:15–3:55

ON SCREEN: The "**Connect ROGUE to your IDE**" section. Click **Add to Cursor**. Cut to Cursor / Claude Desktop: type "What new attacks broke our customer support config in the last 24 hours?" — show it calling `query_diff` + `query_breaches_for_config` and answering with cited results.

 VOICEOVER

"Here's the part that makes ROGUE different. ROGUE doesn't just *consume* Bright Data's MCP server — **it exposes its own**. One click, and your entire threat database is live inside Claude, Cursor, or Windsurf. Ask, in plain English, *what broke our support bot in the last day* — and your AI answers from real breach data, with sources. The red-team intelligence lives where your engineers already work."

8 Business value + close

3:55–4:35

ON SCREEN: Open `/brief` — the CISO-readable daily threat brief, today's diff vs yesterday, Markdown + JSON export buttons. End on the hero, logo, live URL.

 VOICEOVER

"And every morning it ships *this*: a threat brief that's true *today*, not last quarter — what's newly breaking, what you've newly defended, ready to forward to your CISO as Markdown, JSON, or a Slack alert. Every enterprise that ships an LLM has this exposure and no continuous answer to it. ROGUE is that answer — fully deployed, querying the live web, right now."

 VOICEOVER (FINAL BEAT, ON HERO)

"ROGUE. **The red-team that never sleeps.**"

END CARD: `rogue-eosin.vercel.app` · Built solo on all 5 Bright Data products.

DELIVERY NOTES

- **Record the screen live** — the word-rotator, count-up stats, matrix hover-glow, and the MCP query doing real work *are* the demo. Motion is the proof it's alive.

- **Open in incognito** so the first-visit IntroOverlay auto-plays — a built-in cinematic cold open for Scene 1.
- **Pre-warm the API** ~60s before recording (Render free tier cold-starts). Load `/matrix` once first so the heatmap is hot.
- **The two shots you cannot skip:** clicking a red matrix cell → the payload/judge drawer, and the live MCP query in Cursor. They win *Application of Technology* and *Originality* outright. If anything runs long, cut from Scene 6 — never 5 or 7.
- **Judging-criteria map:** Application of Technology → Scenes 3–5 · Originality → Scenes 6–7 · Business Value → Scene 8 · Presentation → the whole cinematic flow.

IF YOU NEED A 60-SECOND CUT

Scene 1 (hook) → Scene 2 hero stats (10s) → Scene 5 matrix red-cell click (20s) → Scene 7 MCP query in Cursor (20s) → Scene 8 tagline. The irreducible "this is real, this is novel" story.
