


Allocation, not technique quality, gates repertoire growth.

A candidate is grey because it was never evaluated — not because it failed. Make it reachable and it graduates.

The least-trying 20 candidates (the tail greedy never reaches)

Greedy

cost-optimal, stops at first breach




candidate quota
suppresses early-stop

forces every candidate to run

Growth

starvation order + quota



 graduated
 evaluated,
no breach
 never
evaluated

reaches ≈ 0 of 20

candidate tier starved by early-stop

8 of 20 graduate (40%)

the other 12 ran but did not breach

The matched A/B (Table 1):

same 40 primitives, same inputs — only the scheduler policy changes.

graduations
0 → **3**

planner reachability
0.07 → **0.98**

starvation
85% → **1%**

And evaluating more is cheaper, not costlier:

cost per graduation \$8.37 at K=3 → \$1.44 at K=20

— the ladder is a fixed cost.