



Containment tracks alignment, not scale

Scale never reduces leakage (within provider)



Alignment contains, at fixed size+family (less aligned → more leakage)



Left: each scale line is WITHIN one provider (Qwen2.5 on Featherless, Llama-3.x on OpenRouter), so its trend is not a serving-stack artifact; open squares are the same model on the other provider (Llama-8B is 83% on Featherless vs 68% on OpenRouter—a 15-pt gap, why we do not pool). Error bars = across-run 95% t-interval (runs=3); 3-32B Qwen saturate at 100% (single run). Right: alignment contrasts within Featherless; abilitated = stock instruct with the refusal direction removed. Two reasoning models (gpt-oss-20b, DeepSeek-R1-70B) leak almost entirely in their chain-of-thought (+87, +97 points).