LIVE · DEPLOYED · OPEN SOURCE

ROGUE

The red-team that never sleeps.

Open-web LLM Threat Intelligence Agent

> **The first two-way Bright Data MCP integration.** The agent harvests the live web through MCP — and exposes its own MCP server so Claude / Cursor users query the live breach matrix from their IDE.

▶ **rogue-eosin.vercel.app**

Soren Obounou Nguia · Web Data UNLOCKED Hackathon (Bright Data × lablab.ai) · May 2026

Powered by Bright Data

02 — PROBLEM

New LLM attacks ship daily. Defenses are tested on yesterday's.

Jailbreaks, prompt-injection patterns, and tool-use exploits surface on Reddit, X, arXiv, GitHub, and security blogs at a pace no static suite tracks. Static red-team tools run fixed payloads against bare models — so a CISO learns about today's attack only after it has already worked on their actual deployment.

Recent in-the-wild jailbreaks — and how fast they spread:

Afrikaans multi-turn	Pliny L1B3RT4S drop	Opus-4.8 tweet jailbreak	Weight-abliteration	Indirect prompt-inject
GPT · Claude · DeepSeek	40+ orgs, repo root	26 full-breach trials in panel	OBLITERATUS HF org	doc-embedded payloads

The window between a jailbreak being posted and it breaching a production chatbot is now **minutes** — not the weeks a quarterly red-team assumes.

03 — THE GAP

Every existing tool sits in the wrong quadrant.

Lakera, HiddenLayer, Promptfoo, Robust Intelligence run fixed suites against bare models. Nobody continuously harvests novel attacks from the open web and reproduces them against your deployment configuration.

Continuous harvest

19 open-web sources, fetched daily — not a frozen test bank.

Your deployment, not a bare model

System prompt + declared tools + target model, scored together.

A closed loop, shipped daily

Harvest → reproduce → diff → brief, every day, to Slack / MCP / dashboard.

	BARE MODEL	YOUR CONFIG
STATIC	Static suites, bare models	Config-aware, but frozen
CONTINUOUS	Continuous, bare models	**ROGUE** continuous × your config

04 — HOW IT WORKS

One five-layer pipeline, end to end.

Harvest → Extract → Dedupe → Reproduce → Diff

1

HARVEST

19 open-web sources via 5 Bright Data products. A bandit tunes SERP spend.

→

2

EXTRACT

An LLM agent structures each document into a portable AttackPrimitive.

→

3

DEDUPE

pgvector cosine similarity clusters near-duplicate attacks.

→

4

REPRODUCE

Each primitive runs against your config × 6 target models × 5 trials.

→

5

DIFF

A judge model verdicts each trial; daily diff ships to Slack / MCP / dashboard.

Bright Data does double duty: it discovers the attacks, and sources the real images multimodal attacks are tested against.

05 — LIVE TODAY

Harvested live — not a fixture.

LIVE FEED

Rows stream in as attacks are harvested — Reddit, X, arXiv, GitHub, MITRE.

BREACH MATRIX

15 attack families × 6 deployments, each cell with a 95% confidence band.

THREAT BRIEF

CISO-readable daily diff — markdown, JSON, Slack, and queryable over MCP.

2 NEW CRITICAL

From one live sweep across 6 deployment configs:

8,321	358	16.5×	44.4%
breach trials judged	attack primitives harvested	spread, most vs least robust	breach rate, PAIR refinement

06 — BRIGHT DATA INTEGRATION

First project to use Bright Data MCP on both sides — harvest and distribution.

CONSUMER the DiscoveryAgent reasons over BD MCP tools to harvest the live web. **PRODUCER** one-click Add to Cursor / VS Code — query the live breach matrix in-IDE, zero setup.

Web Scraper API

Reddit · X · HuggingFace, structured JSON

Social fetches with 0 lines of anti-bot code.

SERP API

Novel-attack discovery, Google + Bing

Self-tuning bandit allocates spend per $.

Web Unlocker

arXiv · blogs · MITRE · OWASP

Clean Markdown straight to the extractor.

Scraping Browser

JS-heavy fallback sites

Reaches archives no pre-built scraper covers.

MCP Server

Agent's primary tool surface

Reasons over tools, no per-product SDK.

$0.15

Bright Data spend per detected breach

$92.48 total BD spend · ~2 min from ingestion to verified breach · self-tuning bandit, $0.05–$0.30 per harvest

Built for the CISO deploying LLMs into production.

Customer-support AI lead

Owns the chatbot customers actually talk to. Needs to know the moment a new attack breaks their setup.

Security / AI red-team lead

Replaces a quarterly manual red-team with an always-on, config-aware feed.

RAG-system owner

Documents the AI reads are an injection surface. ROGUE tests that surface continuously.

MODEL

SaaS threat feed + per-config auto-reproduction tier + enterprise on-prem.

MARKET

LLM security tooling — a fast-growing slice of the broader AI-security spend.

USP

The only continuous-harvest, customer-config-aware closed loop.

Unit economics already proven: **$0.15** Bright Data / breach · **$0.45** all-in / breach · whole 6-day build ran on **$92.48** Bright Data + **$184.22** LLM spend.

08 — WHAT'S NEXT

From hackathon build to deployed feed.

NEXT 30 DAYS

Deeper Web Scraper API coverage — the next 100 sources.

NEXT 90 DAYS

Customer SDK: pipe ROGUE verdicts into SOAR / SIEM.

SCALE

Bright Data Startup Program accelerates Year-1 infrastructure.

> *“I built ROGUE solo in 6 days because Bright Data abstracted away 5 different anti-bot stacks I'd otherwise have spent weeks on. The MCP Server plus pre-built Reddit / X scrapers turned a 6-week project into a 6-day project. Without that infrastructure, this is not possible.”*
>
> — Soren Obounou Nguia · GPTFuzz Grand Prize (Yonsei, 2024) · adversarial-ML research at AIM Intelligence

TRY IT LIVE

rogue-eosin.vercel.app

REPO

`github.com/nguiaSoren/ROGUE`

MCP ENDPOINT

`rogue-api-mr5w.onrender.com/mcp`